

```
In [1]: #importing our required libraries
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
warnings.filterwarnings(action='ignore', category=UserWarning, module='gensim')
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn import metrics
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.neighbors import KNeighborsClassifier
import gensim
import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import TruncatedSVD
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.preprocessing import LabelBinarizer
```

```
In [2]: DB_CONNECT = sqlite3.connect('database.sqlite')
REVIEW_FILTERED = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """,
DB_CONNECT)
# here, we are omitting reviews =3 and then returning new dataframe into REVIEW_FIL
TERED
def polarity(rev):
    if rev<3:
        return 'negative'
    return 'positive'
actualScore=REVIEW_FILTERED['Score'] #score column is copied to actual score
positiveNegative=actualScore.map(polarity)# mapping the column to polarity method a
nd returning strings(positive,negative)
REVIEW_FILTERED['Score']=positiveNegative # replacing the columns in our data with
positive/negative strings
print(REVIEW_FILTERED.shape) #looking at the number of attributes and size of the d
ata
REVIEW_FILTERED.head(2)

(525814, 10)
```

Out [2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0

## TEXT-PRE PROCESSING PHASE

```
In [3]: #remove duplicate rows as we have duplicate product id with similar other data
rem_dup=REVIEW_FILTERED.drop_duplicates(subset={"UserId","ProfileName","Time","Text"},keep="first")
rem_dup.head(2)
```

Out[3]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0

```
In [4]: print(REVIEW_FILTERED.shape)
print(rem_dup.shape)
```

```
(525814, 10)
(364173, 10)
```

after removing duplicates , we have reduced the no .of rows from 525814 to 364173

```
In [5]: # Take out the rows with heplful_num>helpful_den. As W've observed two rows in previous case
df1=rem_dup[rem_dup.HelpfulnessNumerator<=rem_dup.HelpfulnessDenominator]
df1.shape
```

Out[5]: (364171, 10)

```
In [6]: #It's better to sample our data upto 50k points as our dataset is huge
sample_data=5000
df1=df1.sample(sample_data)
df1.sort_values('Time',inplace=True) #Sorting as we want according to time series
df1.head(3)
```

Out [6]:

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>Helpfulness</b>
<b>212494</b>	230308	B00004RYGX	A3C3BAQDZWH5YE	Kushana no shinryaku (Kushana's invasion)	0	1
<b>156021</b>	169214	B0001ES9FI	A2WE1DWPQRDBFW	Scotbuff	6	6
<b>212527</b>	230342	B00004RYGX	A19JYLHD94K94D	Itamar Katz	2	2

```

In [7]: nltk.download('stopwords')
#pre-process each and every text review
'''Function for HTML-Tag removal
we are creating function which removes HTML scripts that present in our text review
'''
def rem_html_tags(review):
    text=re.compile('<.*?>') # regular expressions that need to be find
    return text.sub('',review) #replace nothing while above tags occur
#Function for Punctuation marks removal
def rem_punctuations(review):
    text=re.compile(r'[?|!|\\"|\'|#|.|.|,|)|(|\\|/|~|%|*|]')
    return text.sub('',review)
#Now, Intialise stop-words (most-common words which makes of no-value while vectorin
g)
stop_words=set(stopwords.words('english')) #these are set of stopwords
#Intialise snow-ball stemmer
snow_stem=nltk.stem.SnowballStemmer('english') #returns root of some words
from time import time
%time
count=0 # intializing counter for iterating through entire rows
str1=''
final_str=[]
pos_words=[] #stores words from positive reviews
neg_words=[] #stores words from negative reviews
str2=''
t0=time()
for sentence in df1['Text'].values: # Iterate through all reviews
    filtered_sentence=[] # creating a list which we append filtered words later
    html_filtered=rem_html_tags(sentence) # returns HTML-filtered format
    punct_filtered=rem_punctuations(html_filtered) # returns words with no-punctuation
    marks
    half_filtered=punct_filtered.split() # contains words without HTML & punctuation
    marks
    for word in half_filtered: # iterate through each word in review
        if word.isalpha() and len(word)>2: # checking whether all characters are alpha
            bets and limiting length
            if word.lower() not in stop_words:
                str2=(snow_stem.stem(word.lower())).encode('utf8') # stemming each and ever
                y word
                filtered_sentence.append(str2) # updating our list with stemmed words
                if (df1['Score'].values)[count]=='positive':
                    pos_words.append(str2) #list of all words used to describe positive review
                ws
                if (df1['Score'].values)[count]=='negative':
                    neg_words.append(str2) #list of all words used to describe negative review
                ws reviews
            else:
                continue
            else:
                continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    final_str.append(str1)
    count+=1
print("text successfully pre-processed")

[nltk_data] Downloading package stopwords to
[nltk_data] /home/kprk_25/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
CPU times: user 0 ns, sys: 0 ns, total: 0 ns
Wall time: 5.96 µs
text successfully pre-processed

```

```
In [8]: df1['CleanedText']=final_str # adding new column to df1 by adding our filtered reviews
df1['CleanedText']=df1['CleanedText'].str.decode("utf-8") # standard decoding format
df1.head(2)
```

Out [8]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulness
<b>212494</b>	230308	B00004RYGX	A3C3BAQDZWH5YE	Kushana no shinryaku (Kushana's invasion)	0	1
<b>156021</b>	169214	B0001ES9FI	A2WE1DWPQRDBFW	Scotbuff	6	6

we've added extra column "CleanedText" with pre-processed review

## BAG - OF - WORDS VECTORIZATION

```
In [9]: X_train,X_test,y_train,y_test = train_test_split(df1['CleanedText'].values,df1['Score'].values,test_size=0.3,shuffle=False) # splitting into test and train
# we are not using random state because we've already split based on time
c_vector=CountVectorizer() # initialising bow vectorizer
X_train=c_vector.fit_transform(X_train) # feature extraction of X_train
X_train=preprocessing.normalize(X_train) # normalise data
print("Train Data Size: ",X_train.shape)
X_test=c_vector.transform(X_test) # feature extraction of X_test
X_test=preprocessing.normalize(X_test) # normalise data
print("Test Data Size: ",X_test.shape)
```

```
Train Data Size: (3500, 10650)
Test Data Size: (1500, 10650)
```

we've successfully split data into 70:30 ratio between train and test datasets

```


In [10]: # FINDING BEST K VALUE WITHOUT BRUTE-FORCE OR KD-TREE
from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
'''Now , it is required to perform cross-validation for getting the optimized k-value
firstly, we would go through simple cross-validation on our dataset (cross-val on t
rain data without split)'''
# split the train data set into cross validation train and cross validation test
X_tr,X_cv,y_tr,y_cv = train_test_split(X_train, y_train, test_size=0.3,shuffle=False)
k_list=[]
acc_list=[]
for i in range(1,50,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i)
    k_list.append(i)
    # fitting the model on crossvalidation train
    knn.fit(X_tr, y_tr)
    # predict the response on the crossvalidation train
    pred = knn.predict(X_cv)
    # evaluate CV accuracy
    acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
    acc_list.append(acc)
knn = KNeighborsClassifier(1)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)

```

```

In [11]: err=[]
for i in acc_list:
    err.append(100-i)
sns.set_style("darkgrid")
plt.plot(np.arange(1,50,2),err)
plt.xlabel("K value")
plt.ylabel("error")
plt.title("cross-validation error")
plt.show()
print("minimum error",min(err),'occurs at k =',k_list[acc_list.index(max(acc_list))
])

```



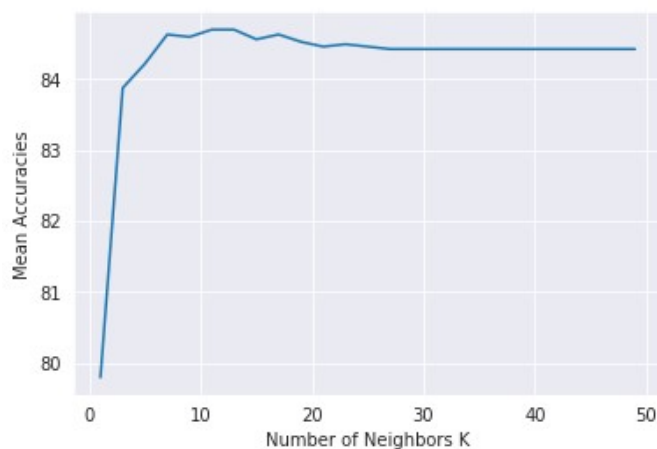
minimum error 15.333333333333329 occurs at k = 9

## Brute-force function

```
In [12]: from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
import time
start_time = time.time()
def func_bruteforce(XTRAIN,YTRAIN):
    knn=KNeighborsClassifier(algorithm='brute')#since, we are using bruteforce algo
    rithm
    grid={'n_neighbors':np.arange(1,50,2)}#creating a dictionary with values as ran
    ge of K-values
    splits=TimeSeriesSplit(n_splits=5)#time series split of train data using 10-fol
    d cross-validation
    grid_search=GridSearchCV(estimator=knn,param_grid=grid,cv=splits,return_train_s
    core=1,n_jobs=-1) #this actually performs grid search over our grid-values as cv va
    lue=10
    grid_search.fit(XTRAIN,YTRAIN)
    # print(grid_search.cv_results_['mean_test_score'])
    return grid_search.best_params_,grid_search.best_score_*100,grid_search.cv_resu
    lts_['mean_test_score']*100# best_params_ and best_score_ are attributes that retur
    ns best k-value and accuracy
    optimal_k,accuracy,acc_scores=func_bruteforce(X_train,y_train)
    print('optimal k :',optimal_k,' with accuracy :',accuracy)
    print('\nlist of accuracies :',acc_scores)
    plt.plot(list(range(1,50,2)),acc_scores)
    plt.xlabel('Number of Neighbors K')
    plt.ylabel('Mean Accuracies')
    plt.show()
    print("time taken: %s seconds" %(time.time()-start_time))
```

optimal k : {'n\_neighbors': 11} with accuracy : 84.69982847341338

list of accuracies : [79.7941681 83.87650086 84.21955403 84.63121784 84.5969125  
2 84.69982847 84.5626072 84.63121784 84.52830189 84.45969125 84.49399657  
84.45969125 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593]



time taken: 22.339608192443848 seconds



## Test-accuracy function

```
In [13]: # NOW, COMPUTE THE TEST-ACCURACY
# instantiate learning model k = optimal_k
def test_acc(optimal_k,XTRAIN,XTEST):
    knn_optimal=KNeighborsClassifier(n_neighbors=optimal_k)
    # fitting the model
    knn_optimal.fit(XTRAIN, y_train)
    # predict the response
    pred=knn_optimal.predict(XTEST)
    # evaluate accuracy
    acc = accuracy_score(y_test, pred) * 100
    return optimal_k,acc,pred
optimal_k,acc,pred=test_acc(optimal_k['n_neighbors'],X_train,X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

The accuracy of the knn classifier for k = 11 is 83.733333%
```

## confusion-matrix function

```

In [14]: from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.preprocessing import LabelBinarizer

def filter(y_test):
    lst=[]
    for x in y_test: # as confusion matrices need only no of positives(1) and negatives(0) in the form of binary digits, we convert labels to +1 or 0
        if x=='positive':
            lst.append(1)
        else:
            lst.append(0)
    return lst


def confuse_matrix(optimal_k,pred,y_test):
    y_con_test=filter(y_test)
    y_con_pred=filter(pred)
    accuracy=accuracy_score(y_con_test,y_con_pred)*100
    print('accuracy we got:', accuracy)
    precision=precision_score(y_con_test,y_con_pred)
    print('precision value :', precision)
    recall=recall_score(y_con_test,y_con_pred)
    print('recall :',recall)
    F1=f1_score(y_con_test,y_con_pred)
    print('F1-SCORE :',F1)
    conf=pd.DataFrame(confusion_matrix(y_con_test,y_con_pred), range(2),range(2))
    sns.set(font_scale=1.4)#for label size
    sns.heatmap(conf, annot=True,annot_kws={"size": 16}, fmt='g') # y-axis-actual,x
    -axis-predicted
    print("      [TN      FP] \n      [FN      TP]")
    return None
confuse_matrix(optimal_k,pred,y_test)

```

```

accuracy we got: 83.73333333333333
precision value : 0.8379287155346334
recall : 0.9975980784627703
F1-SCORE : 0.9108187134502924

```



truncated-svd function

```
In [15]: # for K-D tree, we should pass dence matrce to algorithm, hence, we use Truncated
          SVD
          def TSVD(XTRAIN,XTEST):
              standardized_train = StandardScaler(with_mean=False).fit_transform(XTRAIN) # data
              standardization
              standardized_test = StandardScaler(with_mean=False).fit_transform(XTEST)
              # standardized_train.shape
              tsvd_train = TruncatedSVD(n_components=400, random_state=None).fit_transform(standardized_train)
              # tsvd_train.shape
              tsvd_test = TruncatedSVD(n_components=400, random_state=None).fit_transform(standardized_test)
              # tsvd_test.shape
              return tsvd_train,tsvd_test
          # print(TSVD(X_train,X_test))
```

## KD-Tree Function



```
In [17]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], train_svd, test_svd)
print('\n\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

The accuracy of the knn classifier for k = 7 is 82.933333%
```

## confusion-matrix(kd-tree)

```
In [18]: print(confuse_matrix(optimal_k, pred, y_test))

accuracy we got: 82.93333333333334
precision value : 0.8329979879275654
recall : 0.9943955164131305
F1-SCORE : 0.9065693430656934
      [TN      FP]
      [FN      TP]
None
```



## TF-IDF Vectorization

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(df1['CleanedText'].values, df1['Score'].values, test_size=0.3, shuffle=False) #shuffle=false because we've done time-based split
tfidf = TfidfVectorizer() # initializing vectorizer
X_train = tfidf.fit_transform(X_train) # fitting our model to tfidf
X_train = preprocessing.normalize(X_train) # Normalize Data
print("Train Data Size: ", X_train.shape)
X_test = tfidf.transform(X_test)
X_test = preprocessing.normalize(X_test) # Normalize Data
print("Test Data Size: ", X_test.shape)


Train Data Size: (3500, 10650)
Test Data Size: (1500, 10650)
```

## Brute-Force(TF-IDF)

```
In [20]: start_time = time.time()
optimal_k, accuracy, acc_scores = func_bruteforce(X_train, y_train)
print('optimal k :', optimal_k, ' with accuracy :', accuracy)
print('\nlist of accuracies :', acc_scores)
plt.plot(list(range(1, 50, 2)), acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" % (time.time() - start_time))
```

optimal k : {'n\_neighbors': 9} with accuracy : 84.69982847341338

list of accuracies : [80.17152659 83.77358491 84.5626072 84.45969125 84.69982847 84.5626072 84.5626072 84.45969125 84.49399657 84.45969125 84.45969125 84.45969125 84.45969125 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593]



time taken: 21.356317281723022 seconds

## Test-accuracy on tf-idf

```
In [21]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], X_train, X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 9 is 83.933333%

## confusion matrix on tf-idf (brute)

```
In [22]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 83.93333333333334  
precision value : 0.839622641509434  
recall : 0.9975980784627703  
F1-SCORE : 0.9118185144529821
```

```
[TN      FP]  
[FN      TP]
```

None




## KD-Tree on tf-idf

```
In [23]: start_time = time.time()
train_svd,test_svd=TSVD(X_train,X_test) # passing our vectorized array to TSVD function
optimal_k,accuracy,acc_scores=func_kdtree(train_svd,y_train)
print('optimal k :',optimal_k,' with accuracy :',accuracy)
print('\nlist of accuracies :',acc_scores)
plt.plot(list(range(1,50,2)),acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" %(time.time()-start_time))
```

optimal k : {'n\_neighbors': 5} with accuracy : 84.6655231560892

list of accuracies : [80.92624357 84.08233276 84.66552316 84.5626072 84.5283018  
9 84.59691252  
84.52830189 84.49399657 84.49399657 84.45969125 84.49399657 84.45969125  
84.45969125 84.45969125 84.45969125 84.45969125 84.45969125 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593]



time taken: 348.4332015514374 seconds

## test-acc on tf-idf(kd-tree)

```
In [24]: optimal_k,acc,pred=test_acc(optimal_k['n_neighbors'],train_svd,test_svd)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 5 is 82.666667%

## confusion-matrix on tf-idf(kd-tree)



```
In [25]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 82.66666666666667
precision value : 0.8338960162052668
recall : 0.988791032826261
F1-SCORE : 0.9047619047619048
      [TN      FP]
      [FN      TP]
None
```



## Word-2-Vector Model

```
In [26]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sent=[] # list used to store the words seperated with commas for building o
ur w2vec model
for sent in df1['CleanedText'].values:
    list_of_sent.append(sent.split())# split acc. to space (default split)
```

```
In [27]: print(df1['CleanedText'].values[0]) # actual text contained in cleaned text
print("*****")
print(list_of_sent[0]) # seperated words
```

```
night big fan cartoon shown decid watch also tri see winona ryder movi know well
known fact like see danc end movi hey funni apart hilari film like much thing wr
ong short mayb special efect bit hey older
*****
['night', 'big', 'fan', 'cartoon', 'shown', 'decid', 'watch', 'also', 'tri', 'se
e', 'winona', 'ryder', 'movi', 'know', 'well', 'known', 'fact', 'like', 'see', '
danc', 'end', 'movi', 'hey', 'funni', 'apart', 'hilari', 'film', 'like', 'much',
'thing', 'wrong', 'short', 'mayb', 'special', 'efect', 'bit', 'hey', 'older']
```

```
In [28]: # min_count = 1 considers all the words
w2v_model=Word2Vec(list_of_sent,min_count=1,size=300, workers=4)
```

```
In [29]: w2v_words = list(w2v_model.wv.vocab)
# print("number of words that occured minimum 1 time ",len(w2v_words))
# # print("sample words ", w2v_words[0:NumOfWords])
```

```
In [30]: w2v_model.wv.most_similar('friend')
```

```
Out[30]: [('away', 0.9998213052749634),  
          ('special', 0.9997816681861877),  
          ('sure', 0.9997462034225464),  
          ('plan', 0.9997340440750122),  
          ('probabl', 0.999731183052063),  
          ('offic', 0.9997212290763855),  
          ('sardin', 0.9997199773788452),  
          ('experi', 0.9997192621231079),  
          ('given', 0.9997184872627258),  
          ('trap', 0.9997171759605408)]
```

```
In [31]: w2v_model.wv.most_similar('milk')
```

```
Out[31]: [('add', 0.9994142055511475),  
          ('coconut', 0.999239981174469),  
          ('hot', 0.9991186857223511),  
          ('ad', 0.998521089553833),  
          ('bitter', 0.9979516267776489),  
          ('mix', 0.9979391694068909),  
          ('dark', 0.9978399276733398),  
          ('sugar', 0.9977403879165649),  
          ('strong', 0.9967581033706665),  
          ('ice', 0.9966917037963867)]
```

## Avg- word 2 vector

```
In [32]: # average Word2Vec
# compute average word2vec for each review.
final_list=[]; # list which stores avg-w2v for each review
for sent in list_of_sent: # iterating through each review
    sent_vec = np.zeros(300) # intializing an array to 300 dimensions to store all
    avg-w2vec in document
    count=0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word] # converting word to vector at this instance
            sent_vec+= vec # adding each vector upto our length of document
            count+=1 # increment-word count
    if count!=0: # base condition that satisfies for each document (min count of words >=1)
        sent_vec/=count #avg of all word-vectors
    final_list.append(sent_vec) # appending avg-w2v to our list
print(len(final_list))
print(len(final_list[0]))
print(final_list[0])
```

```
5000
300
[-0.01164844 -0.16398207 -0.1598202 -0.20830556 0.1930184 -0.00523259
 0.12709514 0.1379925 0.31362519 -0.12860666 -0.12486024 0.11189036
-0.19413995 -0.14901359 0.00896853 -0.05037179 -0.11900527 0.03789675
-0.08068039 0.30614107 0.33141025 0.1074885 -0.14101324 0.032429
 0.21435868 0.07603103 0.31960051 -0.01160748 0.01687335 0.13656699
 0.06671033 0.06207677 0.16056856 -0.08613736 0.2423953 0.0169799
 0.28325437 0.02052751 -0.16824489 -0.06350347 -0.00683602 0.28731411
 0.06820311 0.1762026 -0.00340028 0.1096158 0.12631391 -0.17586311
 0.30316582 -0.16262445 -0.00711274 0.18104894 -0.20172886 -0.38927122
 0.21718782 -0.02526711 0.04045336 -0.29768907 0.02897178 0.03854182
 0.40598411 0.09192314 0.24544751 0.32983695 -0.01685455 -0.04935611
 0.00132356 -0.24400696 0.03461177 -0.27146973 0.20008281 -0.14294173
-0.03765911 0.00714741 0.00833894 0.07264583 0.00798375 0.03662911
-0.0110697 -0.10806979 0.07848909 0.10159594 -0.34940116 0.26006845
-0.07386274 0.13974425 -0.00866134 -0.05865972 -0.02103276 -0.00301911
 0.01385339 -0.14276742 -0.08659495 -0.04270958 0.23170534 0.07741351
 0.05222722 -0.19070475 0.18224015 -0.07949573 0.36034243 0.00155192
 0.00830711 -0.14419432 -0.05177989 -0.12778855 -0.10901433 0.18364578
 0.05057086 -0.13506567 -0.21614412 -0.28822348 -0.11845554 0.14554334
 0.04860891 -0.16780589 -0.04104192 -0.04644938 -0.18320662 -0.25848072
 0.07599313 -0.10301752 0.14396914 -0.2149244 0.06868274 0.28940474
 0.02876104 -0.18202684 -0.02380706 0.06590119 0.0938291 -0.0446062
 0.18656738 0.37455977 0.19908624 -0.15168581 0.10127564 -0.1817282
 0.00514612 -0.170241 -0.06648799 -0.2349549 -0.03528513 -0.14385264
-0.28933547 0.18944508 0.05109583 -0.20412021 -0.00804618 0.01080682
-0.25059177 0.23617235 0.19567608 0.37867181 0.15982852 0.08623538
-0.19596971 0.07748037 0.11503273 0.05292407 0.10599103 0.10117855
-0.09859948 0.08470739 0.09460978 0.32814257 -0.05210048 -0.24181528
-0.23487202 0.11476227 0.15061621 -0.18705424 0.08762189 0.17893813
 0.02946271 -0.04831937 0.20117948 0.17530106 0.10373197 -0.03205493
 0.06611548 -0.00960309 -0.09490015 0.15123264 -0.21414927 0.14745754
 0.04951481 -0.05316664 -0.11658236 0.23171285 -0.01026776 0.3486508
 0.01565099 0.06060475 0.23756456 0.11461154 -0.20055038 -0.27191999
 0.15725437 -0.23014316 -0.20362457 0.09232868 -0.24040721 -0.03563088
-0.00355863 0.33249253 0.09047104 -0.01805996 -0.14880588 -0.12776537
-0.34148613 0.08886645 -0.25846928 0.26399506 -0.09887185 0.093373
 0.09862953 0.49247571 0.13216332 0.27079449 -0.0760092 0.01815752
 0.04323808 -0.32273575 0.02779741 -0.09735397 -0.18735652 0.17102059
-0.08231306 0.06877979 0.34730217 -0.17193846 0.14805185 0.14608719
-0.26245513 0.40121633 -0.14494346 0.11200341 -0.23264397 0.20673764
-0.02229906 0.2180251 -0.11826002 0.02305463 0.43900355 0.19570877
 0.40081416 -0.22487193 0.29947373 0.09516151 -0.01459126 0.04712402
 0.5499556 0.12535784 -0.05457577 0.1957064 -0.03101439 -0.04273791
-0.25045779 -0.23135917 0.45719389 -0.16452721 0.03729769 0.21248902
 0.04346795 0.31566773 -0.06288796 -0.2338725 -0.20585889 0.16074302
 0.2764035 0.12315805 -0.02243407 -0.01451979 0.05088301 0.38277082
 0.0036519 0.01732661 -0.12970312 -0.02524693 0.09920473 0.05906569
 0.19291627 0.32661416 0.29847595 -0.12735559 -0.16375066 -0.28276342
 0.16382955 0.02126821 0.14005213 0.22682169 0.21535663 0.00441957
 0.23248557 -0.1285177 0.11835978 -0.04353898 0.01795444 0.06512267]
```

```
In [33]: np.isnan(final_list).any() # checking if any undefined values(NaN) present in our data
```

```
Out[33]: False
```


```
In [34]: # final_list_norm=preprocessing.normalize(final_list)
X_train,X_test,y_train,y_test=train_test_split(final_list,df1['Score'].values,test_
size=0.3,shuffle=False)
# time basis split, hence, shuffle =false
#X_train[0]
final_list=np.asarray(final_list)
```

## brute-force on avg- w2 vec

```
In [35]: start_time = time.time()
optimal_k,accuracy,acc_scores=func_bruteforce(X_train,y_train)
print('optimal k :',optimal_k,' with accuracy :',accuracy)
print('\nlist of accuracies :',acc_scores)
plt.plot(list(range(1,50,2)),acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" %(time.time()-start_time))

optimal k : {'n_neighbors': 17}  with accuracy : 84.4253859348199

list of accuracies : [76.432247    81.23499142  83.25900515  83.97941681  84.2195540
3  84.28816467
 84.28816467  84.39108062  84.42538593  84.39108062  84.3567753   84.39108062
 84.39108062  84.42538593  84.42538593  84.42538593  84.42538593  84.42538593
 84.42538593  84.42538593  84.42538593  84.42538593  84.42538593  84.42538593
 84.42538593]
```



time taken: 40.699716567993164 seconds

## test-accuracy on avg-w2v(brute)


```
In [36]: optimal_k,acc,pred=test_acc(optimal_k['n_neighbors'],X_train,X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

The accuracy of the knn classifier for k = 17 is 83.266667%
```

## confusion-matrix on avg-w2v(brute)

```
In [37]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 83.26666666666667  
precision value : 0.8326666666666667  
recall : 1.0  
F1-SCORE : 0.9086940705711167  
      [TN      FP]  
      [FN      TP]  
None
```



## K-D tree on avg-w2vec

```
In [38]: start_time = time.time()
# AS AVG-W2V IS dense matrix , we no need to apply truncated SVD on this model
optimal_k, accuracy, acc_scores = func_kdtree(X_train, y_train)
print('optimal k :', optimal_k, ' with accuracy :', accuracy)
print('\nlist of accuracies :', acc_scores)
plt.plot(list(range(1, 50, 2)), acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" % (time.time() - start_time))

optimal k : {'n_neighbors': 17}  with accuracy : 84.4253859348199

list of accuracies : [76.432247    81.23499142  83.25900515  83.97941681  84.2195540
 3 84.28816467
 84.28816467 84.39108062 84.42538593 84.39108062 84.3567753  84.39108062
 84.39108062 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593
 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593
 84.42538593]
```

```
time taken: 70.69144010543823 seconds
```

### test-acc on avg w2v(KD)


```
In [39]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], X_train, X_test)
print('\n\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

The accuracy of the knn classifier for k = 17 is 83.266667%
```

### confusion-matrix on avg-w2v(KD-TREE)

```
In [40]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 83.26666666666667
precision value : 0.8326666666666667
recall : 1.0
F1-SCORE : 0.9086940705711167
      [TN      FP]
      [FN      TP]
None
```



## TF-IDF WEIGHTED W2VEC

Take weighted sum of the vectors divided by the sum of all the tfidf's ----->  $(\text{tfidf}(\text{word}) \times \text{w2v}(\text{word})) / \text{sum}(\text{tfidf's})$

----->another type of vectorization technique to convert sentence into vector notation



```

In [41]: # X_train,X_test,y_train,y_test=train_test_split(df1['CleanedText'],df1['Score'].values,
          # test_size=0.3,shuffle=False)
          final_tf_idf=tfidf.fit_transform(df1['CleanedText'].values)
          print(final_tf_idf.shape)# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
          # TF-IDF weighted Word2Vec
          features= tfidf.get_feature_names() # tfidf words/col-names
          tfidf_final_list=[]; # the tfidf-w2v for each sentence/review is stored in this list
          rows=0;
          for sent in list_of_sent: # for each review/sentence
              sent_vec=np.zeros(300) # intializing an array to 300 dimensions to store all tf-idf-w2vec in document
              weight_sum=0; # num of words with a valid vector in the sentence/review
              for word in sent: # iterating through each review
                  if word in w2v_words:
                      vec=w2v_model.wv[word] # converting word to vector at this instance
                      tf_idf=final_tf_idf[rows,features.index(word)] # obtain the tf_idf of a word in a sentence/review
                      sent_vec+=(vec*tf_idf) #multiplying tf-idf value of word with w2v vector of that word
                      weight_sum+=tf_idf # summing tf_idf value of every word
                  if weight_sum!= 0: # base condition that satisfies for each document(min weighted sum >=1)
                      sent_vec/= weight_sum # for each document ,this step converts sentence to tf-idf-w2v format by dividing TF-IDF sentence vectors with sum of all tf-idf's
                      tfidf_final_list.append(sent_vec)
                      rows+=1

          (5000, 13166)

```

```

In [42]: type(final_tf_idf) #

```

```

Out[42]: scipy.sparse.csr.csr_matrix

```

```

In [43]: tfidf_final_list=np.asarray(tfidf_final_list)
          X_train, X_test, y_train, y_test=train_test_split(tfidf_final_list,df1['Score'].values,
          test_size=0.3,shuffle=False)
          tfidf_final_list.shape

```

```

Out[43]: (5000, 300)

```

```

In [44]: np.isnan(tfidf_final_list).any() # checking if any undefined values (NaN) present in our data

```

```

Out[44]: False


```

## brute-force on TF-IDF w2v

```
In [45]: start_time = time.time()
optimal_k, accuracy, acc_scores = func_bruteforce(X_train, y_train)
print('optimal k :', optimal_k, ' with accuracy :', accuracy)
print('\nlist of accuracies :', acc_scores)
plt.plot(list(range(1, 50, 2)), acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" % (time.time() - start_time))
```

optimal k : {'n\_neighbors': 15} with accuracy : 84.4253859348199

list of accuracies : [77.11835334 81.09777015 83.3619211 83.97941681 84.3567753  
84.28816467  
84.3567753 84.42538593 84.42538593 84.39108062 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593]



time taken: 26.69974637031555 seconds

## Test-acc on tf-idf w2v (brute)

```
In [46]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], X_train, X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```


The accuracy of the knn classifier for k = 15 is 83.066667%

## confusion matrix on tf-idf w2v(brute)

```
In [47]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 83.06666666666666  
precision value : 0.8323313293253173  
recall : 0.9975980784627703  
F1-SCORE : 0.9075018208302986
```

```
      [TN      FP]  
      [FN      TP]  
None
```

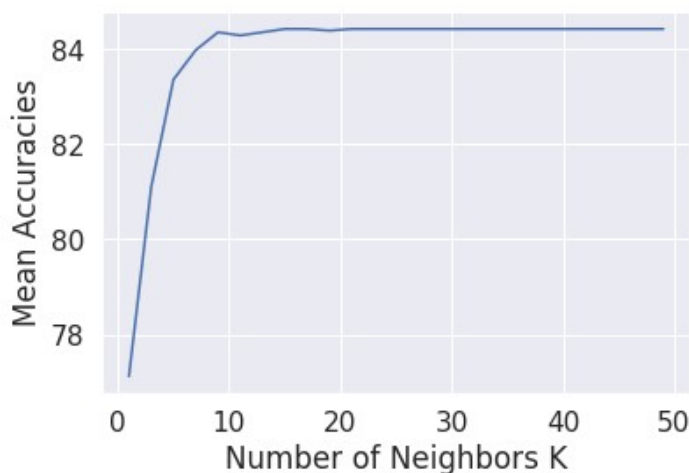


**kd-tree on tf-idf w2v (brute)**

```
In [48]: start_time = time.time()
# AS tf-idf w2v IS dense matrix , we no need to apply truncated SVD on this model
optimal_k, accuracy, acc_scores = func_kdtree(X_train, y_train)
print('optimal k :', optimal_k, ' with accuracy :', accuracy)
print('\nlist of accuracies :', acc_scores)
plt.plot(list(range(1, 50, 2)), acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" % (time.time() - start_time))
```

optimal k : {'n\_neighbors': 15} with accuracy : 84.4253859348199

list of accuracies : [77.11835334 81.09777015 83.3619211 83.97941681 84.3567753  
84.28816467  
84.3567753 84.42538593 84.42538593 84.39108062 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593]



time taken: 52.86696982383728 seconds

## test-acc on tfidf w2v (kd-tree)

```
In [49]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], X_train, X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 15 is 83.066667%

## confusion-matrix on tfidf-w2v (kd-tree)

```
In [50]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 83.06666666666666
precision value : 0.8323313293253173
recall : 0.9975980784627703
F1-SCORE : 0.9075018208302986
      [TN      FP]
      [FN      TP]
None
```



## BAG - OF - WORDS (Bi-grams)

we've done vectorization on uni-grams using bag-of-words. Bi-gram model will take all unique sets of adjacent two words and then count their occurrences to assign frequency

```
In [52]: X_train,X_test,y_train,y_test = train_test_split(df1['CleanedText'].values,df1['Score'].values,test_size=0.3,shuffle=False) # splitting into test and train
# we are not using random state because we've already split based on time
c_vector=CountVectorizer(ngram_range=(1,2)) # initialising bow vectorizer of bi-grams
X_train=c_vector.fit_transform(X_train) # feature extraction of X_train
X_train=preprocessing.normalize(X_train) # normalise data
print("Train Data Size: ",X_train.shape)
X_test=c_vector.transform(X_test) # feature extraction of X_test
X_test=preprocessing.normalize(X_test) # normalise data
print("Test Data Size: ",X_test.shape)
```


```
Train Data Size: (3500, 107862)
Test Data Size: (1500, 107862)
```

## Brute-force algorithm(bi-grams-BOW)

```
In [53]: start_time = time.time()
optimal_k, accuracy, acc_scores = func_bruteforce(X_train, y_train)
print('optimal k :', optimal_k, ' with accuracy :', accuracy)
print('\nlist of accuracies :', acc_scores)
plt.plot(list(range(1, 50, 2)), acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" % (time.time() - start_time))
```

optimal k : {'n\_neighbors': 5} with accuracy : 84.6655231560892

list of accuracies : [80.24013722 84.11663808 84.66552316 84.66552316 84.5969125  
2 84.5626072  
84.5626072 84.49399657 84.45969125 84.45969125 84.42538593 84.45969125  
84.45969125 84.45969125 84.45969125 84.45969125 84.45969125 84.45969125  
84.45969125 84.45969125 84.45969125 84.45969125 84.42538593 84.42538593  
84.42538593]



time taken: 1945.3550109863281 seconds

## test-acc on bruteforce(bi-gram)

```
In [54]: optimal_k, acc, pred = test_acc(optimal_k['n_neighbors'], X_train, X_test)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 5 is 84.466667%

## confusion matrix (bi-grams-bruteforce)

```
In [56]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 84.46666666666667  
precision value : 0.8522884882108183  
recall : 0.9839871897518014  
F1-SCORE : 0.9134150873281308
```

```
[TN      FP]  
[FN      TP]
```

None

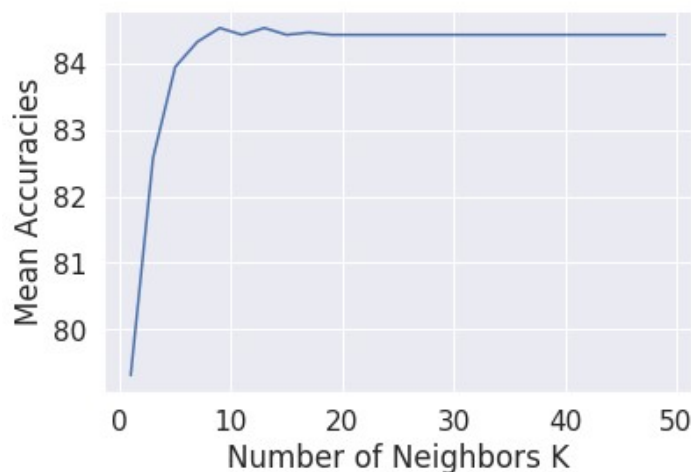


## KD-TREE(bi-gram)

```
In [58]: start_time = time.time()
train_svd,test_svd=TSVD(X_train,X_test) # passing our vectorized array to TSVD function
optimal_k,accuracy,acc_scores=func_kdtree(train_svd,y_train)
print('optimal k :',optimal_k,' with accuracy :',accuracy)
print('\nlist of accuracies :',acc_scores)
plt.plot(list(range(1,50,2)),acc_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Mean Accuracies')
plt.show()
print("time taken: %s seconds" %(time.time()-start_time))
```

optimal k : {'n\_neighbors': 9} with accuracy : 84.52830188679246

list of accuracies : [79.31389365 82.5728988 83.94511149 84.32246998 84.5283018  
9 84.42538593  
84.52830189 84.42538593 84.45969125 84.42538593 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593 84.42538593 84.42538593 84.42538593 84.42538593 84.42538593  
84.42538593]



time taken: 418.3834536075592 seconds

## test-acc on kd-tree(bi-grams)

```
In [59]: optimal_k,acc,pred=test_acc(optimal_k['n_neighbors'],train_svd,test_svd)
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 9 is 82.933333%

## conf-matrix(kd-tree) on bi-grams



```
In [60]: print(confuse_matrix(optimal_k,pred,y_test))
```

```
accuracy we got: 82.93333333333334  
precision value : 0.8321070234113712  
recall : 0.9959967974379503  
F1-SCORE : 0.9067055393586007  
      [TN      FP]  
      [FN      TP]  
None
```



## observations(Accuracies on each vectorization):

BOW(uni-gram) BOW(bi-gram) TF-IDF AVG-W2V TF-IDF-W2V Brute-force 83.73% 84.46% 83.93% 83.26% 83.06% KD-TREE  
82.93% 82.93% 82.66% 83.26% 83.06%

## CONCLUSIONS

- >As observed, Highest accuracy was achieved by bag-of-words(using bi-grams), 84.46%
- >For Avg-w2v and TF-IDF w2v , BRUTE-FORCE and KD-TREE yielding similar accuracies
- >CV accuracy and test-accuracy doesn't changes much for same algorithm
- >taking a lot of time for training(with truncated svd, it takes much more time to train)